

# Arithmetic Solving in Z3

Nikolaj Bjørner and Lev Nachmanson

Microsoft

**Abstract.** The theory of arithmetic is integral to many uses of SMT solvers. Z3 has implemented native solvers for arithmetic reasoning since its first release. We present a full re-implementation of Z3's original arithmetic solver. It is based on substantial experiences from user feedback, engineering and experimentation. While providing a comprehensive overview of the main components we emphasize selected new insights we arrived at while developing and testing the solver.

## 1 Introduction

The theory of arithmetic is among the most prolific theories used in SMT solvers. It is used across a wide set of applications, and they have a wide range of demands. Supporting efficient theory solvers for arithmetic requires balancing feature support, ranging from linear, difference logic, to linear real, linear integer, non-linear polynomial arithmetic and in cases transcendental functions such as exponentiation. The aim of this paper is to provide a high-level, yet self-contained, overview of internal ingredients of the arithmetic solver. It seeks to explain tool users what to expect of solving methodologies when using Z3 for arithmetic. We assume familiarity of basics of SMT solving using CDCL(T), e.g., [5]. While make an effort to cover all features for completeness, we devote attention to highlight a selection that to our knowledge are unique for arithmetic solvers. These highlights include (1) how the new solver patches linear real programming solutions to find solutions to integer variables, (2) heuristics that are new in how the solver finds Gomory cuts, and (3) the solver's integration of Gröbner basis computation for solving non-linear constraints. User pain points around SMT have to our experience centered dominantly around quantifiers and non-linear arithmetic. A complete (for non-linear reals) solver that integrates with other theories and quantifier reasoning becomes relevant. To evaluate the contribution of each feature we use benchmarks drawn from SMTLIB benchmark sets [3] and benchmarks supplied by a user, Certora [12].

In overview, the arithmetic solver uses a waterfall model for solving arithmetic constraints. It is illustrated with additional details in Figure 1.

- First, it establishes feasibility with respect to linear inequalities. Variables are solved over the rationals; Section 3.
- Second, it establishes feasibility with respect to mixed integer linear constraints; Integer variables are considered solved if they are assigned integer values while satisfying linear inequalities; Section 4.
- Finally, it establishes feasibility with respect to non-linear polynomial constraints; Section 5.

The rest of the paper elaborates on the components of the solver. For completeness, we go through all relevant pieces. We highlight parts that to our knowledge are novel.

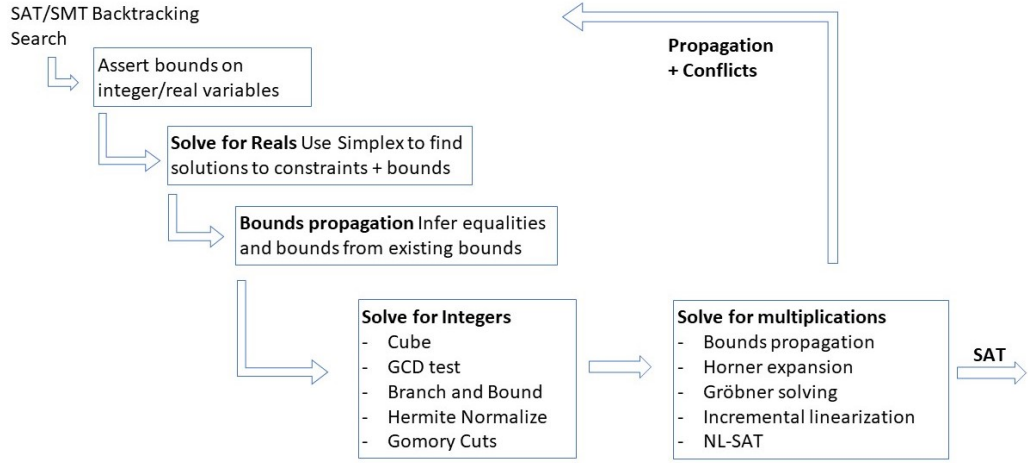


Fig. 1: Overview of Z3's Arithmetic Theory Solver

## 2 Design Goals and Implementation Choices

The SMT formalism for arithmetic in many cases subsumes formalisms used by mixed integer, MIP, solvers. However, there are several fundamental differences between the workloads we have tuned the arithmetic solver for compared to workloads seen by MIP solvers. Z3 uses infinite precision “big-num” numeral representations, in contrast to using floating points. The drawback is that the arithmetic solver is impractical on linear programming optimization problems, but the engine avoids having to compensate for rounding errors and numerical instability. The solver uses a sparse matrix representation for the Dual Simplex tableau. We also created a version that uses an LRU decomposition and floating point numbers but found that extending this version with efficient backtracking was not practical compared to the straight-forward sparse matrix format. Finally, the solver remains integrated within a CDCL engine that favors an eager case split strategy leaving it to conflict analysis to block infeasible branches. This contrasts mainstream MIP designs that favor a search tree of relatively few branches where the engine performs significant analysis before case splits.

## 3 Linear Real Arithmetic

The solver first determines whether arithmetic constraints are feasibility over the reals. It also attempts to propagate equalities eagerly for shared variables and infer stronger bounds of variables.

### 3.1 Linear Solving

Based on [21] the solver for real linear inequalities uses a dual simplex solver. It partitions the variables into *basic* and *non-basic* variables, and maintains a global set of equalities of the form  $x_{bi} + \sum_j a_{ij}x_j = 0$ , where  $i$  refers to the  $i$ 'th row,  $x_{bi}$  is basic and  $x_j$  range over non-basic variables. It also maintains an evaluation  $\beta$ , such that  $\beta(x_{bi}) + \sum_j a_{ij}\beta(x_j) = 0$  for each row  $i$ . Each variable  $x_j$  is assigned lower and upper bounds during search. The solver then checks

whether  $lo_j \leq \beta(x_j) \leq hi_j$ , for bounds  $lo_j, hi_j$  that are dynamically added and removed by Boolean decisions  $x_j \leq hi_j$ ,  $x_j \geq lo_j$ . If the bounds are violated, it updates the evaluation and pivots if necessary. We recall the approach using an example.

*Example 1.* For the following formula

$$y \geq 0 \wedge (x + y \leq 2 \vee x + 2y \geq 6) \wedge (x + y \geq 2 \vee x + 2y > 4)$$

the solver introduces auxiliary variables  $s_1, s_2$  and represents the formula as

$$x + y - s_1 = 0, x + 2y - s_2 = 0, x \geq 0, (s_1 \leq 2 \vee s_2 \geq 6), (s_1 \geq 2 \vee s_2 > 4)$$

Only bounds (e.g.,  $s_1 \leq 2$ ) are asserted during search. The slack variables  $s_1, s_2$  are initially basic (dependent) and  $x, y$  are non-basic. In dual Simplex tableaux, values of a non-basic variable  $x_j$  can be chosen between  $lo_j$  and  $hi_j$ . The value of a basic variable is a function of non-basic variable values. Pivoting swaps basic and non-basic variables and moves basic variables within their bounds to bounds violations. For example, assume we start with a set of initial values  $x = y = s_1 = s_2 = 0$  and bounds  $x \geq 0, s_1 \leq 2, s_1 \geq 2$ . Then  $s_1$  has to be 2 and it is made non-basic. Instead,  $y$  becomes basic:  $y + x - s_1 = 0, s_2 + x - 2s_1 = 0$ . The new tableau updates the assignment of variables to  $x = 0, s_1 = 2, s_2 = 4, y = 2$ . The resulting assignment is a model for the original formula.

### 3.2 Finding equal variables - cheaply

It is useful to have the arithmetic solver propagate implied equalities when arithmetic is used in combination with other theories, or even when it solves non-linear arithmetic constraints. Equality propagation is disabled for pure arithmetic theories, such as `QF_LIA`, `QF_LRA` [3]. Z3 originally used a method based on storing *offset* equalities in a hash table. An offset equality is of the form  $x_i = y + k$ , where  $k$  is a numeric constant. Offset equalities are extracted from rows that contain  $x_i$  as a basic variable, and contains only one other non-fixed variable  $y$ , while other variables are fixed and their lower (upper) bounds add up to  $k$ . It turns out that computing  $k$  is expensive when the tableau contains large numeric constants. Hash table operations contribute with additional overhead. It turns out that neither the offset hash-table, nor computing  $k$ , is really necessary. We describe our new method, using an example. We first described the method for avoiding to compute offsets in [8]. The description there relies on building a tree data-structure for connecting variables and fails to leverage that the dual simplex tableau can be used directly.

*Example 2.* From equalities  $x + 1 = y, y - 1 = z$  infer that  $x = z$ . Based on the tableau form, the solver is presented with the original equality atoms via slack variables

$$s_1 = x + u - y, s_2 = y - u - z, 1 \leq u \leq 1$$

The tableau can be solved by setting  $x = 0, y = 0, z = 0, s_1 = 1, s_2 = -1, u = 1$ . The slack variables are bounded when the equalities are asserted

$$s_1 = x + u - y, s_2 = y - u - z, 1 \leq u \leq 1, 0 \leq s_1 \leq 0, 0 \leq s_2 \leq 0$$

The original solution is no longer valid, the values for  $s_1, s_2$  are out of bounds. Pivoting re-establishes feasibility using a different solution, for example

$$x = z - u - s_1, y = z - u - s_2, 1 \leq u \leq 1, 0 \leq s_1 \leq 0, 0 \leq s_2 \leq 0$$

with assignment  $z = 0, x = y = -1$ . The variables  $x$  and  $y$  have the same value, but must they be equal under all assignments? We can establish this directly by subtracting the right-hand sides  $z - u - s_1$  and  $z - u - s_2$  from another and by factoring in the constant bounds to obtain the result 0. But subtraction is generally expensive if there are many bounded constants in the rows. Such arithmetical operations are not required to infer that  $x = y$ .

Z3 uses the following conditions to infer an equality between variables  $x, y$  having the same values in the current assignment:

- $x$  is basic, and the tableau has row  $x - y + \alpha = 0$ ,
- $x, y$  are connected through a non-basic variable  $z$  in a pair of the tableau rows in one of the following forms (1)  $x - z + \alpha = 0, y - z + \alpha' = 0$ , (2)  $x + z + \alpha = 0, y + z + \alpha' = 0$ ,

where  $\alpha, \alpha'$  are linear combinations of fixed variables.

We experimented with generalizing the connection between equal variables to allow non-unit coefficients on  $z$ , but it did not result in measurable improvements.

### 3.3 Bounds Propagation

It is not uncommon that SMT formulas contain different bounds for the same variable, such as one atom  $x \geq 2$  and another atom  $x \geq 3$ . When the atom  $x \geq 3$  is assigned to true, the solver can directly propagate  $x \geq 3$ . Bounds can also be inferred indirectly. With a row  $x - 2y = 0$  and bound  $y \geq 1$ , it follows that  $x \geq 2$ . To implement direct bounds propagation, the solver maintains an index that maps each variable to the set of bounds atoms where it occurs. To implement indirect bounds propagation, the solver queries updated rows for whether they imply bounds that are stronger than the currently asserted bounds. If so, these stronger bounds are used by the index for direct bounds propagation.

## 4 Integer Linear Arithmetic

The mixed integer linear solver consists of several layers that first attempt to *patch* integer variables from solutions over reals to solutions over integers. Then, if patching fails to correct all integer variables, it checks for integer infeasibility by checking light-weight Diophantine feasibility criteria and then resort to variants of Gomory Cuts and Branch and Bound.

### 4.1 Patching

In a feasible tableau we can assume that all non-basic variables are at their bounds and therefore if they have integer sort they are assigned integer values. Only the basic variables could be assigned non-integer values. Patching seeks changing values of non-basic values in order to assign integer values to basic variables. A related method, that diversifies values of variables using *freedom* intervals was described in [18], but we found it does not preserve integral assignments.

Thus, we patch rows with basic variables  $x_b \notin \mathcal{Z}$ . We use a process that seeks a  $\delta$ , such that  $|\delta|$  is minimal and the row with  $x_b$  is of the form  $x_b + \alpha y + \alpha' x' = 0$ , where  $\alpha \notin \mathcal{Z}$ , such that the update  $\beta(y) := \beta(y) + \delta$  is within the bounds of  $y$ ,  $x_b$  is assigned an integer value and such that  $x_b$  becomes integer without breaking any bounds in the tableau. We describe details of the patching method in Appendix A.1.

*Example 3.* Suppose we are given a tableau of the form  $y - \frac{1}{2}x = 0$ ,  $z - \frac{1}{3}x = 0$  where  $y, z$  are basic variables and  $x$  has bounds  $[3, 10]$ ,  $y$  has bounds  $[-3, 4]$ ,  $z$  has bounds  $[-5, 12]$ . The

variable  $x$  is initially assigned at the bound  $\beta(x) = 3$ . Then  $\beta(y) = \frac{3}{2}$  and  $\beta(z) = 1$ . But neither  $y$  nor  $z$  is close to their bounds. We can move  $x$  to 8 without violating the bound for  $y$  because of  $y - \frac{1}{2}x = 0$ . Thus, the freedom interval for  $x$  is the range  $[3, 8]$  and within this range there is a solution,  $x = 6$ , where  $y, z$  are integers and within their bounds.

## 4.2 Cubes

An important factor in solving more satisfiable integer arithmetic instances is a method by Bromberger and Weidenbach [9,10]. It allows detecting feasible inequalities over integer variables by solving a stronger linear system. Their method relies on the following property: The inequalities  $Ax \leq b$  are integer feasible, for matrix  $A$  and vectors  $x, b$ , if  $Ax \leq b - \frac{1}{2}\|A\|_1$  has a solution over the reals. We use the 1-norm  $\|A\|_1$  of a matrix as a column vector, such that each entry  $i$  is the sum of the absolute values of the elements in the corresponding row  $A_i$ .

*Example 4.* Suppose we have  $3x + y \leq 9 \wedge -3y \leq -2$  and wish to find an integer solution. By solving  $3x + y \leq 9 - \frac{1}{2}(3 + 1) = 7, -3y \leq -2 - \frac{1}{2}3 = -3.5$  we find a model where  $y = \frac{7}{6}, x = 0$ . After rounding  $y$  to 1 and maintaining  $x$  at 0 we obtain an integer solution to the original inequalities.

Z3 includes a twist relative to [10] that allows to avoid strengthening on selected inequalities [6]. First, we note that *difference* inequalities of the form  $x - y \leq k$ , where  $x, y$  are integer variables and  $k$  is an integer offset need not be strengthened: they have a solution over reals if and only if they have a solution over integers. For octagon constraints  $\pm x \pm y \leq k$ , there is a boundary condition: they need only require strengthening if  $x, y$  are assigned at mid-points between integral solutions. For example, if  $\beta(x) = \frac{1}{2}$  and  $\beta(y) = \frac{3}{2}$ , for  $x + y \leq 2$ .

## 4.3 GCD consistency

A basic test for integer infeasibility is by enforcing divisibility constraints.

*Example 5.* Assume we are given a row  $5/6x + 3/6y + z + 5/6u = 0$ , where  $x, y$  are fixed at  $2 \leq x \leq 2, -1 \leq u \leq -1$ , and  $z$  is the base variable. Then it follows that  $5 + 3(y + 2z) = 0$  which has no solution over the integers: The greatest common divisor of coefficients to the non-fixed variables (3) does not divide the constant offset from the fixed variables (5).

The basic test is extended as follows. For each row  $ax + by + c = 0$ , where

- $a, b, c$  and  $x, y$  are vectors of integer constants and variables, respectively.
- the coefficients in  $a$  are all the same and smaller than the coefficients in  $b$
- the variables  $x$  are bounded

Let  $l := a \cdot lb(x), u := a \cdot ub(x)$ . That is, the lower and upper bounds for  $a \cdot x$  based on the bounds for  $x$ . If  $\lfloor \frac{u}{\gcd(b,c)} \rfloor > \lceil \frac{l}{\gcd(b,c)} \rceil$ , then there is no solution for  $x$  within the bounds for  $x$ .

## 4.4 Branching

Similar to traditional MIP branch-and-bound methods, the solver creates somewhat eagerly case splits on bounds of integer variables if the dual simplex solver fails to assign them integer values. For example, Simplex may assign an integer variable  $x$ , the value  $\frac{1}{2}$ , in which case z3 creates a literal  $x \leq 0$  that triggers two branches  $x \leq 0$  and  $\neg(x \leq 0) \equiv x \geq 1$ .

## 4.5 Cuts

The arithmetic solver produces Gomory cuts from rows where the basic variables are non-integers after the non-basic variables have been pushed to the bounds. Z3 implements Chvátal-Gomory cuts described in [21]. It also implements algorithms from [20,13] to generate cuts after the linear systems have been transformed into Hermitian matrices. It is a long-standing and timely challenge [1] to harness the effectiveness of selecting cuts. While the solver takes [21] as starting point, it incorporates a few heuristics and enhancements.

Recall that a row  $\sum_{j=0}^k a_j \cdot x_j + x_b = 0$  from the tableau is called a Gomory row, and is eligible for Gomory cut, if  $x_b$  is a basic variable and  $x_j$  are non-basic variables,  $x_b$  is an integral variable, but  $\beta(x_b)$  is not integral, and for each  $x_j$  we have  $\beta(x_j) = lo_j$  or  $\beta(x_j) = hi_j$ , and the bounds are not strict.

We use a relaxed definition of Gomory rows. For a non-basic integral variable  $x_j$  we allow for value  $\beta(x_j)$  to be not at a bound when  $\beta(x_j)$ , and  $a_j$  are both integers: Let us call such  $x_j$  *row integral*. We provide further justification for this relaxed definition in Appendix A.2.

To select a cut variable, our main heuristic sorts all Gomory rows from the tableau by the distance of  $\beta(x_b)$  from the nearest integer, that is  $\min\{\beta(x_b) - \lfloor \beta(x_b) \rfloor, \lceil \beta(x_b) \rceil - \beta(x_b)\}$ , and pick a few of them having the minimal distance to produce the cuts. We break the ties by preferring the variables that are used in more terms. Heuristics used previously relied on distances to bounds.

We also look for the case when  $x_b$  is at an extremum. For example, if for all  $x_j$  we have  $\beta(x_j) = lo_j$ , and  $a_j > 0$  then  $x_b$  is at the maximum, and we deduce  $x_b \leq \lfloor \beta(x_b) \rfloor$ . The explanations of the Gomory term do not include constraints on  $x_j$  for  $j \in A$  from the relaxed definition, but in case of an extremum these constraints should be added.

Cuts are consequences of the current bounds. By default the solver adds new rows to the Dual Simplex tableau corresponding to cuts, but makes an exception when the new rows include large numerals. In analogy the solver avoids bounds propagation, Section 3.3, when computation of bounds relies on big-num arithmetic. Similarly, cuts that involve large coefficients are first added to a temporary scope where the tableau is checked for feasibility. The cuts are only re-added within the main scope if the temporary tableau is infeasible.

## 5 Non-Linear Arithmetic

Similar to solving for integer feasibility, the arithmetic solver solves constraints over polynomials using a waterfall model for non-linear constraints. At the basis it maintains, for every monomial term  $x \cdot x \cdot y$ , a definition  $m = x \cdot x \cdot y$ , where  $m$  is a variable that represents the monomial  $x \cdot x \cdot y$ . The module for non-linear arithmetic then collects the monomial definitions that are violated by the current evaluation, that is  $\beta(m) \neq \beta(x) \cdot \beta(x) \cdot \beta(y)$ . It attempts to establish a valuation  $\beta'$  where  $\beta'(m) = \beta'(x) \cdot \beta'(x) \cdot \beta'(y)$ , or derive a consequence that no such evaluation exists.

### 5.1 Patch Monomials

A *patch* for a variable  $x$  is *admissible* if the update  $\beta(x) := v$  does not break any integer linear constraints and  $x$  does not occur in monomial equations that are not already false under  $\beta$ .

- Set  $\beta(m) := \beta(x) \cdot \beta(x) \cdot \beta(y)$  and check if the patch of  $m$  is admissible.
- Try to set  $\beta(y) := \beta(m) / (\beta(x) \cdot \beta(x))$ , provided  $\beta(x)$  is not 0, and check that the patch for  $x$  is admissible.
- When  $\beta(m) = r^2$  for a rational and  $m := x \cdot x$  try patching  $x$  by setting  $\beta(x) := \pm r$ .

## 5.2 Bounds propagation

A relatively inexpensive step is to propagate and check bounds based on non-linear constraints. For example, for  $y \geq 3$ , then  $m = x \cdot x \cdot y \geq 3$ , if furthermore  $x \leq -2$ , we have the strengthened bound  $m \geq 12$ . Bounds propagation can also flow from bounds on  $m$  to bounds on the variables that make up the monomial, such that when  $m \geq 8, 1 \leq y \leq 2, x \leq 0$ , then we learn the stronger bound  $x \leq -2$  on  $x$ . It uses an interval arithmetic abstraction, that understands bounds propagation over squares. Thus, if  $-2 \leq x \leq 2$ , then  $0 \leq x^2 \leq 4$  instead of  $-4 \leq x^2 \leq 4$ .

The solver also performs Horner expansions of polynomials to derive stronger bounds. For example, if  $x \geq 2, y \geq -1, z \geq 2$ , then  $y + z \geq 1$  and therefore  $x \cdot (y + z) \geq 2$ , but we would not be able to deduce this fact if combining bounds individually for  $x \cdot y$  and  $x \cdot z$  because no bounds can be inferred for  $x \cdot y$  in isolation. The solver therefore attempts different re-distribution of multiplication in an effort to find stronger bounds.

## 5.3 Adding bounds

Non-linear bounds propagation only triggers if all variables are either bounded from above or below or occur with an even power. The solver includes a pass where it adds a bound  $x \geq 0$  to variables  $x$  where  $lo_x = -\infty, hi_x = +\infty$ . The added bound may help trigger bounds propagation, such as detecting conflicts on  $xy > 0, xz > 0, y > 0 > z$ .

## 5.4 Gröbner reduction

Z3 uses a best effort Gröbner basis reduction to find inconsistencies, cheaply, and propagate consequences. While Gröbner basis heuristics are not new to Z3, they have evolved and to our knowledge the integration is unique among SMT solvers. Recall that reduced Gröbner basis for a set of polynomial equations  $p_1 = 0, \dots, p_k = 0$  is a set  $q_1 = 0, \dots, q_m = 0$ , such that every  $p_i$  is a linear sum of  $q_j$ 's, and the leading monomials of every pair  $q_i, q_j, i \neq j$ , have no common factors. Since Z3 uses completion as a heuristic to make partial inferences, it does not seek to compute a basis. The Gröbner module performs a set of partial completion steps, preferring to eliminate variables that can be isolated, and expanding a bounded number of super-position steps (reductions by S-polynomials).

Z3 first adds equations  $m = x_1 \dots x_k$  for monomial definitions that are violated. It then traverses the transitive cone of influence of Simplex rows that contain one of the added variables from monomial definitions. It only considers rows where the basic variable is bounded. Rows where the basic variable is unbounded are skipped because the basic variable can be solved for over the reals. Fixed variables are replaced by constants, and the bounds constraints that fixes the variables are recorded as dependencies with the added equation. Thus, the equations handled by the Gröbner basis reduction are of the form  $\langle p_i : xy + 3z + 3 = 0, d_i : \{3 \leq u \leq 3\} \rangle$ , where  $p_i$  is a polynomial and  $d_i$  is a set of dependencies corresponding to fixed variables that were replaced by constants in  $p_i$ . In the example, we replaced  $u$  by 3 and the definition  $\langle m = xy, \emptyset \rangle$  resolved  $m$  by  $xy$ . Dependencies are accumulated when two polynomials are resolved to infer a new derived equality. Generally, when  $\langle xy + p_1 = 0, d_1 \rangle, \langle xz + p_2 = 0, d_2 \rangle$  are two polynomial equations, then  $\langle zp_1 - yp_2 = 0, d_1 \cup d_2 \rangle$  can be derived accumulating the premises  $d_1, d_2$ .

Finally, equations are pre-solved if they are linear and can be split into two groups, one containing a single variable that has a lower (upper) bound, the other with more than two variables with upper (lower) bounds. This avoids losing bounds information during completion.

After (partial) completion, the derived equations are post-processed:

**Constant propagation** For equalities of the form  $x = 0$  or  $ax + b = 0$ . If the current assignment to  $x$  does not satisfy the equation, then the equality is propagated as a lemma.

**Linear propagation** As a generalization of constant propagation, if the completion contains linear equations that evaluate to false under the current assignment, then these linear equations are added to the Simplex Tableau. Example 6 illustrates a use where this propagation is useful.

**Factorization** Identify factors of the form  $xyp \simeq 0$  where  $x, y$  are variables and  $p$  is linear. We infer the clause  $xyp \simeq 0 \Rightarrow x \simeq 0 \vee y \simeq 0 \vee p \simeq 0$ .

*Example 6 (Combining Gröbner completion and Linear Solving).* We include an example obtained from Yoav Rodeh at Certora. The instance was not solvable prior to adding simplex propagation. To solve it, Certora relied on treating multiplication as an uninterpreted function and including selected axioms for modular arithmetic and multiplication that were instantiated by E-matching. The distilled example is:

$$L \leq x \cdot y \leq U \wedge 1 \leq x \wedge m_r \leq U \wedge x \cdot y \neq m_r$$

where  $L = N \text{ div } 2, U = 1 + L, m_r = (x \cdot (\text{ite}(y \geq 0, y, N + y))) \bmod N$ . We assume  $N$  is even, such as  $N = 2^{256}$ . The solver associates a variable  $m$  with  $x \cdot y$  and  $m'$  with  $x \cdot y'$  and  $y'$  with  $\text{ite}(y \geq 0, y, N + y)$  and includes the constraints  $0 \leq m_r < N, m_q \cdot N + m_r = m'$ , where  $m_q$  is an integer variable. The most interesting case is where  $y < 0$ , so  $y' = y + N$ . Gröbner basis completion then allows to derive  $m_q N + m_r = m' = x(y + N) = xy + xN = m + xN$ , which by integer linear arithmetic reasoning (the extended GCD test) contradicts  $m \neq m_r$  because the absolute value of both variables is below  $N$ .

Our extraction of linear constraints represents a partial integration of linear programming and polynomial arithmetic, that favors only including linear inequalities over variables and monomials that are already present. Our implementation does not include any variables for new monomials produced by completion. In comparison, the approach in [25] proposes a domain for abstract interpretation that populates a linear solver with all equations produced by a completion. We have not experimented in depth with extending our approach with a full basis, or use it as a starting point for finding lemmas based on Positivstellensatz or other extension mechanisms [34,32].

We use an adaptation of ZDD (Zero suppressed decision diagrams [29,30]) to represent polynomials. The representation has the advantage that polynomials are stored in a shared data-structure and operations over polynomials are memorized. A polynomial over the real is represented as an acyclic graph, where nodes are labeled by variables and edges are labeled by coefficients. Figure 2 shows a polynomial stored in a polynomial decision diagram, PDD.

The root node labeled by  $x$  represents the polynomial  $x \cdot l + r$ , where  $l$  is the polynomial of the left sub-graph and  $r$  the polynomial of the right sub-graph. The left sub-graph is allowed to be labeled again by  $x$ , but the right sub-graph may only have nodes labeled by variables that are smaller in a fixed ordering. The fixed ordering used in this example sets  $x$  above  $y$ . Then the polynomial for the right sub-graph is  $y + 1$ , and the polynomial with the left sub-graph is  $5xy + (y + 1)$ .




Fig. 2: PDD representation of  $5x^2y + xy + y + x + 1$



## 5.5 Incremental linearization

Following [14] we incrementally linearize monomial definitions that currently evaluate to false. For example, we include lemmas of the form  $x = 0 \rightarrow m = 0$  and  $x = 1 \rightarrow m = y$ , for  $m = x^2y$ . Incremental linearization proceeds by first applying linearizations that are considered cheap, such as case splitting on whether variables take values 0, 1, -1, when these boundary conditions are exhausted, instantiates lemmas based on monotonicity of multiplication and tangents. It is possible that there are overlapping monomial definitions, such as  $m' = x \cdot y$ . Then incremental linearization takes into account that the definition for  $m$  can be *factored* into  $m' \cdot x$ . It also uses specialized congruence closure reasoning, recognizing equalities modulo signs, such that when  $m = x \cdot y$ ,  $m' = z \cdot y$  and  $x = -z$  in the current context, then  $m \sim -m'$ .

To find all factorizations of monomial  $m = \prod_{i \in A} x_i$  as  $m = m_0 \cdot m_1$ , we choose  $a \in A$  and enumerate over all proper subsets  $B$  of  $A$  containing  $a$ . For each  $B$  we check that  $m_0 = \prod_{i \in B} x_i$  and  $m_1 = \prod_{i \in A \setminus B} x_i$  are monomials.

To support floating point arithmetic reasoning we also include incremental linearization lemmas for special cases of exponentiation [15]. We also added rules for incremental linearization of divisibility operations. The front-end to the core arithmetic solver axiomatizes integer and real division operations using multiplication and addition, so that the solver does not have to reason about division. Nevertheless, we found use cases for axioms of the form  $y > 0 \wedge x > z \Rightarrow x/y > z/y$ , bypassing indirect reasoning around constraints created by axioms.

## 5.6 NLSat

As an end-game attempt, the solver attempts to solve the non-linear constraints using a complete solver for Tarski's fragment supported by the NLSat solver [24]. NLSAT is complete for non-linear arithmetic and includes branch-and-bound to handle cases of integer arithmetic. It can therefore sometimes be used to solve goals, bypassing the partial heuristics entirely. The solver therefore includes selected calls to NLSat with a small resource bound to close branches before attempting incomplete heuristics such as incomplete linearization. The results in Section 7 suggests that our use of NLSat with a resource bound currently incurs significant overhead on easy problems, but overall is an advantage. We found that it is sometimes the case that turning off NLSat all-together can speed up the solver significantly, but is overall a disadvantage.

## 6 Shared Equalities

Z3 uses model-based theory combination [18] for sharing equalities between theories. In the context of arithmetic it means that in a satisfiable state shared variables where  $\beta(x) = \beta(y)$  it also holds that the literal  $x \simeq y$  is assigned to true. For larger benchmarks we observed that there can be a significant overhead in checking whether a term occurs in a shared context because it relies on properties of which parent terms it occurs. We therefore introduced a way to cache the property of being shared in the E-nodes. The property gets invalidated when the a new parent E-node is added or the congruence class of the E-node is merged.

## 7 Evaluation

To get an idea of how the new solver compares and how the individual features of the new solver weigh on performance we conducted a set of measurements. They are based on three benchmark sets: QF-LIA, SMTLIB2 benchmarks for the theory of quantifier-free integer linear arithmetic,

QF\_NIA, SMTLIB2 benchmarks for the theory of quantifier-free non-linear integer arithmetic, and **benchmark-submission**, a smaller set of verification conditions obtained from Certora. Data associated with the measurements summarized in this Section are available from [7]. The measurements are listed in Appendix B. We ran the solvers for 600s and measured how many problems are solved within 600s. We compared default settings of the solvers with CVC5 [2,26,17], MathSat5 [16,28], and Yices2 [22,37], and Z3’s legacy arithmetic solver, which is available by setting the option `smt.arith.solver=2`. The advances relative to the legacy solver are noticeable. Compared to other solvers, Yices2 and MathSat5 shine as fast out of the gates solving relatively more problems within 1s, but are mainly limited by the set of features it supports, such as lack of support for algebraic data-types.

The feature-wise evaluation suggests that using NLSat to eagerly close branches comes with a steep cost for easy benchmarks. It can likely be tuned in future versions of Z3. The eager use of NLSat still provides an overall benefit. Z3 also uses *tactics* that run a few strategies with a 5 second resource bounds early on to find models using SAT encodings and selected branch-and-bound strategies. They are also a cause of relatively slow startup. The default tactics can be overridden. The overall biggest impact feature is incremental linearization. While it is run after gcd tests, bounds propagation and Gröbner saturation, it has a significant effect. Other features have each a relative minor effect in isolation. The solver relies on their cumulative effect.

## 8 Summary and Discussion

We presented the architecture and a cross-cut of system innovations in a new arithmetic solver in Z3. It shows to provide good advances relative to the legacy arithmetic solver, and our evaluation suggests it compares very well with other state-of-art SMT solvers. The new solver enabled us to address some design choices with the previous solver that limited extensibility. Notably, the new solver separates its representation of arithmetic constraints from terms shared by other solvers through an E-graph. We noticed that limitation of using shared terms is that the boundary for when to treat a sub-term as a variable or a polynomial is inherently ambiguous. The legacy solver is also highly incomplete for non-linear reasoning (over the reals).

Many avenues for further innovations and tuning remain. Another important aspect is trust. Implementing the many features of the arithmetic solver is inherently a complex task. Many bugs get uncovered by fuzzing [11,27,35,36,31,23,33] both in the legacy and new solver, bearing witness to the difficulty of creating a correct solver. The solver therefore supports a number of ways to validate results. The easiest validation is for satisfiable formulas, where the satisfiable formula is model checked against the returned model. The main difficulty with satisfiable models is to correctly track interpretations of under-specified operations, such as division by 0. To check that consequences produced by the solver are valid, there is a self-validator enabled by the `smt.arith.validate=true`. It uses the legacy arithmetic solver to check lemmas and propagations. There is also a mechanism for creating certificates that can be processed offline or online. With each theory axiom and propagation produced by the solver, it produces a certificate object that can be used to validate inferences by the arithmetic solver. The certificates are exposed in proof objects [19] and also as annotations in proof logs [4]. Z3 contains a built-in proof checker for proof logs. The proof checker for arithmetic certificates validates conflicts that can be justified by using Farkas lemma and bounds propagations that use cuts. It currently falls back to invoking Z3 on lemmas (using the legacy arithmetic solver) for non-linear lemmas and other cases not covered by the built-in checker. Certificates created for QF\_LRA are fully handled, while self-contained or independent proof checking for more expressive fragments of arithmetic is future work.

## References

1. Maria-Florina Balcan, Siddharth Prasad, Tuomas Sandholm, and Ellen Vitercik. Structural analysis of branch-and-cut and the learnability of gomory mixed integer cuts. In Sanmi Koyejo, S. Mohamed, A. Agarwal, Danielle Belgrave, K. Cho, and A. Oh, editors, *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*, 2022.
2. Haniel Barbosa, Clark W. Barrett, Martin Brain, Gereon Kremer, Hanna Lachnitt, Makai Mann, Abdalrhman Mohamed, Mudathir Mohamed, Aina Niemetz, Andres Nötzli, Alex Ozdemir, Mathias Preiner, Andrew Reynolds, Ying Sheng, Cesare Tinelli, and Yoni Zohar. cvc5: A versatile and industrial-strength SMT solver. In Dana Fisman and Grigore Rosu, editors, *Tools and Algorithms for the Construction and Analysis of Systems - 28th International Conference, TACAS 2022, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2022, Munich, Germany, April 2-7, 2022, Proceedings, Part I*, volume 13243 of *Lecture Notes in Computer Science*, pages 415–442. Springer, 2022.
3. Clark Barrett, Pascal Fontaine, and Cesare Tinelli. The Satisfiability Modulo Theories Library (SMT-LIB), 2016.
4. Nikolaj Bjørner. Proofs for SMT. <https://z3prover.github.io/slides/proofs.html>, 2022.
5. Nikolaj Bjørner, Clemens Eisenhofer, Arie Gurfinkel, Nuno P. Lopes, Leonardo de Moura, Lev Nachmanson, and Christoph Wintersteiger. Z3 internals. <https://z3prover.github.io/papers/z3internals.html>, 2023.
6. Nikolaj Bjørner and Lev Nachmanson. Theorem recycling for theorem proving. In *Vampire*, 2017.
7. Nikolaj Bjørner and Lev Nachmanson. Supplementary data. <https://github.com/z3prover/doc/arithmatic>, 2024.
8. Nikolaj S. Bjørner and Lev Nachmanson. Navigating the universe of Z3 theory solvers. In Gustavo Carvalho and Volker Stolz, editors, *Formal Methods: Foundations and Applications - 23rd Brazilian Symposium, SBMF 2020, Ouro Preto, Brazil, November 25-27, 2020, Proceedings*, volume 12475 of *Lecture Notes in Computer Science*, pages 8–24. Springer, 2020.
9. Martin Bromberger and Christoph Weidenbach. Fast cube tests for LIA constraint solving. In *IJCAR*, 2016.
10. Martin Bromberger and Christoph Weidenbach. New techniques for linear arithmetic: cubes and equalities. *Formal Methods in System Design*, 51(3):433–461, 2017.
11. Robert Brummayer, Florian Lonsing, and Armin Biere. Automated testing and debugging of SAT and QBF solvers. In Ofer Strichman and Stefan Szeider, editors, *Theory and Applications of Satisfiability Testing - SAT 2010, 13th International Conference, SAT 2010, Edinburgh, UK, July 11-14, 2010. Proceedings*, volume 6175 of *Lecture Notes in Computer Science*, pages 44–57. Springer, 2010.
12. Certora. Certora Benchmarks. <https://github.com/jar-ben/benchmark-submission>, 2023.
13. Jürgen Christ and Jochen Hoenicke. Cutting the mix. In *CAV*, 2015.
14. Alessandro Cimatti, Alberto Griggio, Ahmed Irfan, Marco Roveri, and Roberto Sebastiani. Experimenting on solving nonlinear integer arithmetic with incremental linearization. In *SAT*, 2018.
15. Alessandro Cimatti, Alberto Griggio, Ahmed Irfan, Marco Roveri, and Roberto Sebastiani. Incremental linearization for satisfiability and verification modulo nonlinear arithmetic and transcendental functions. *ACM Trans. Comput. Log.*, 19(3):19:1–19:52, 2018.
16. Alessandro Cimatti, Alberto Griggio, Bastiaan Joost Schaafsma, and Roberto Sebastiani. The mathsat5 SMT solver. In Nir Piterman and Scott A. Smolka, editors, *Tools and Algorithms for the Construction and Analysis of Systems - 19th International Conference, TACAS 2013, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2013, Rome, Italy, March 16-24, 2013. Proceedings*, volume 7795 of *Lecture Notes in Computer Science*, pages 93–107. Springer, 2013.
17. CVC5. CVC5 executable. [https://cvc5.stanford.edu/downloads/builds/x86\\_64-win64-production/cvc5-2024-01-08-x86\\_64-win64-production.exe](https://cvc5.stanford.edu/downloads/builds/x86_64-win64-production/cvc5-2024-01-08-x86_64-win64-production.exe), 2024.
18. Leonardo Mendonça de Moura and Nikolaj Bjørner. Model-based theory combination. *Electron. Notes Theor. Comput. Sci.*, 198(2):37–49, 2008.

19. Leonardo Mendonça de Moura and Nikolaj Bjørner. Proofs and refutations, and Z3. In Piotr Rudnicki, Geoff Sutcliffe, Boris Konev, Renate A. Schmidt, and Stephan Schulz, editors, *Proceedings of the LPAR 2008 Workshops, Knowledge Exchange: Automated Provers and Proof Assistants, and the 7th International Workshop on the Implementation of Logics, Doha, Qatar, November 22, 2008*, volume 418 of *CEUR Workshop Proceedings*. CEUR-WS.org, 2008.
20. Isil Dillig, Thomas Dillig, and Alex Aiken. Cuts from proofs: A complete and practical technique for solving linear inequalities over integers. In *CAV*, 2009.
21. B. Dutertre and L. de Moura. A Fast Linear-Arithmetic Solver for DPLL(T). In *CAV*, 2006.
22. Bruno Dutertre. Yices 2.2. In Armin Biere and Roderick Bloem, editors, *Computer Aided Verification - 26th International Conference, CAV 2014, Held as Part of the Vienna Summer of Logic, VSL 2014, Vienna, Austria, July 18-22, 2014. Proceedings*, volume 8559 of *Lecture Notes in Computer Science*, pages 737–744. Springer, 2014.
23. Declan Hwang. Z3 Issue tracker. <https://github.com/z3prover/z3>, 2024.
24. Dejan Jovanovic and Leonardo Mendonça de Moura. Solving non-linear arithmetic. In *IJCAR*, 2012.
25. Zachary Kincaid, Nicolas Koh, and Shaowei Zhu. When less is more: Consequence-finding in a weak theory of arithmetic. *Proc. ACM Program. Lang.*, 7(POPL):1275–1307, 2023.
26. Gereon Kremer, Andrew Reynolds, Clark W. Barrett, and Cesare Tinelli. Cooperating techniques for solving nonlinear real arithmetic in the cvc5 SMT solver (system description). In Jasmin Blanchette, Laura Kovács, and Dirk Pattinson, editors, *Automated Reasoning - 11th International Joint Conference, IJCAR 2022, Haifa, Israel, August 8-10, 2022, Proceedings*, volume 13385 of *Lecture Notes in Computer Science*, pages 95–105. Springer, 2022.
27. Muhammad Numair Mansur, Maria Christakis, Valentin Wüstholtz, and Fuyuan Zhang. Detecting critical bugs in SMT solvers using blackbox mutational fuzzing. In Prem Devanbu, Myra B. Cohen, and Thomas Zimmermann, editors, *ESEC/FSE '20: 28th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering, Virtual Event, USA, November 8-13, 2020*, pages 701–712. ACM, 2020.
28. MathSat5. MathSat5 executable. <https://mathsat.fbk.eu/download.php?file=mathsat-5.6.10-win64-msvc.zip>, 2024.
29. Shin-ichi Minato. Zero-suppressed bdds for set manipulation in combinatorial problems. In Alfred E. Dunlop, editor, *DAC*, 1993.
30. Masaaki Nishino, Norihito Yasuda, Shin-ichi Minato, and Masaaki Nagata. Zero-suppressed sentential decision diagrams. In *AAAI*, 2016.
31. Jiwon Park, Dominik Winterer, Chengyu Zhang, and Zhendong Su. Generative type-aware mutation for testing SMT solvers. *Proc. ACM Program. Lang.*, 5(OOPSLA):1–19, 2021.
32. André Platzer, Jan-David Quesel, and Philipp Rümmer. Real world verification. In Renate A. Schmidt, editor, *Automated Deduction - CADE-22, 22nd International Conference on Automated Deduction, Montreal, Canada, August 2-7, 2009. Proceedings*, volume 5663 of *Lecture Notes in Computer Science*, pages 485–501. Springer, 2009.
33. Maolin Sun, Yibiao Yang, Yang Wang, Ming Wen, Haoxiang Jia, and Yuming Zhou. Smt solver validation empowered by large pre-trained language models. In *2023 38th IEEE/ACM International Conference on Automated Software Engineering (ASE)*, pages 1288–1300, 2023.
34. Ashish Tiwari. An algebraic approach for the unsatisfiability of nonlinear constraints. In C.-H. Luke Ong, editor, *Computer Science Logic, 19th International Workshop, CSL 2005, 14th Annual Conference of the EACSL, Oxford, UK, August 22-25, 2005, Proceedings*, volume 3634 of *Lecture Notes in Computer Science*, pages 248–262. Springer, 2005.
35. Dominik Winterer, Chengyu Zhang, and Zhendong Su. On the unusual effectiveness of type-aware operator mutations for testing SMT solvers. *Proc. ACM Program. Lang.*, 4(OOPSLA):193:1–193:25, 2020.
36. Dominik Winterer, Chengyu Zhang, and Zhendong Su. Validating smt solvers via semantic fusion. In *Proceedings of the 41st ACM SIGPLAN Conference on Programming Language Design and Implementation*, pages 718–730, 2020.
37. Yices2. Yices2 executable. [https://yices.csl.sri.com/releases/2.6.4/yices-2.6.4-x86\\_64-pc-mingw32-static-gmp.zip](https://yices.csl.sri.com/releases/2.6.4/yices-2.6.4-x86_64-pc-mingw32-static-gmp.zip), 2024.

## A Integers

### A.1 Patching

We here outline the method for patching in more detail.

Given a row  $x_b + \alpha y + r = 0$ , where  $x_b$  is a basic variable,  $y$  is non-basic variable multiplied by the fraction  $\alpha$ , and  $r$  is the remaining of the row, the shift amount  $\delta$  is computed based on the following analysis.

Take first the fractional parts of  $\beta(x_b)$  and  $\alpha$ :

- $f_x := x_1/x_2 := \text{frac}(\beta(x_b))$ , s.t.  $0 < x_1 < x_2$ , and  $x_1, x_2$  are mutually prime.
- $f_\alpha := a_1/a_2 := \text{frac}(\alpha)$ , s.t.  $0 < a_1 < a_2$ , and  $a_1, a_2$  are mutually prime.

The goal is to compute integers

- $\min \delta^+ > 0$  .  $\text{frac}(\alpha\delta^+) = 1 - f_x$ ,
- $\max \delta^- < 0$  .  $\text{frac}(\alpha\delta^-) = 1 - f_x$ .

These two integers are the minimal amount to move the value  $\beta(y)$  to make the value of  $\beta(x_b)$  integral.

Let  $\mathcal{Z}$  be the set of integers. We solve for  $\delta \in \mathcal{Z}$  such that  $L := \frac{x_1}{x_2} + \frac{a_1}{a_2}\delta \in \mathcal{Z}$  too. If  $L \in \mathcal{Z}$  then  $a_2 \frac{x_1}{x_2} + a_1\delta$  is also an integer. Therefore  $a_2 \frac{x_1}{x_2} \in \mathcal{Z}$ . That means  $a_2 := tx_2$  for some  $t \in \mathcal{Z}$ , because  $x_1$  and  $x_2$  are coprime. By substituting  $a_2$  with  $x_2t$  we get  $L := \frac{x_1}{x_2} + \frac{a_1}{x_2t}\delta$  and  $Lx_2 := x_1 + \frac{a_1}{t}\delta \in \mathcal{Z}$ . Since  $t \uparrow a_2$ , and  $a_2$  and  $a_1$  are coprime,  $t \uparrow \delta$ . Therefore, we search for  $\delta$  in form  $\delta := mt$ ,  $m \in \mathcal{Z}$ . We obtain  $L := \frac{x_1}{x_2} + \frac{a_1}{x_2t}mt = \frac{x_1 + ma_1}{x_2}$ . From  $L \in \mathcal{Z}$  follows  $x_2k = x_1 + ma_1$  for some  $k \in \mathcal{Z}$ . We can rewrite the last equality as  $x_1 = a_1m - x_2k$ . Because  $x_2 \uparrow a_2$ , and  $a_1$  and  $a_2$  are coprime,  $x_2$  and  $a_1$  are mutually prime too. That means that for some  $u, v$  we have  $1 = a_1u + x_2v$ . We can show that if  $\delta := utx_1$  then  $\frac{x_1}{x_2} + \frac{a_1}{a_2}\delta \in \mathcal{Z}$ .

From the other side, for any  $\gamma \in \mathcal{Z}$  satisfying  $\frac{x_1}{x_2} + \frac{a_1}{a_2}\gamma \in \mathcal{Z}$  holds  $\frac{a_1(\delta - \gamma)}{a_2} \in \mathcal{Z}$ . Since  $a_1$  and  $a_2$  are coprime,  $a_2 \uparrow (\delta - \gamma)$ , and  $\gamma := \delta \bmod a_2$ . We conclude that  $\delta^+ = \delta \bmod a_2$ , and  $\delta^- = \delta^+ - a_2$ .

### A.2 Row Integral Variables

To show our reasoning we represent the set of non-basic row indices as the union of two disjoint subsets  $A \cup B$ , where  $\{x_j : j \in A\}$  is the set of all row integral variables and  $B$  is the rest of non-basic indices. The row then can be written as  $\sum_{j \in A} a_j \cdot x_j + \sum_{j \in B} a_j \cdot x_j + x_b = 0$ , that is equivalent to  $\sum_{j \in A} a_j \cdot x_j + x_b = -\sum_{j \in B} a_j \cdot x_j$ .

By choice of  $A$  and the fact that  $x_b$  is integral, the left-hand side has to be integral, but currently the fractional part of the left side value is equal to  $\beta(x_b) - \lfloor \beta(x_b) \rfloor$ , called  $f_0$  in [21]. Further on we can repeat all the steps of the proof of the Gomory inequality from [21], starting from the observation that the right-hand side value change should be either greater than or equal to  $1 - f_0$ , or smaller than or equal to  $-f_0$ , for the left-hand side to become integral.

## B Evaluation Data

We compare how many instances the solvers handle within 1s, within 1-10s, 10-100s, and 10-600s. We also list timeouts and cases where the solver returns unknown because of incompleteness, and cases where benchmarks are unhandled, either because the solver runs out of allocated virtual memory set to 2GB, or due to unsupported features. The version of Z3 used for the experiments corresponds to 4.12.5.

Solver	< 1s	1 to 10s	10 to 100s	100 to 600s	> 600s	unknown/unhandled	solved
CVC5	3082	4564	3578	1693	10959	0/0	12917
MathSat5	3304	6022	3894	2047	8607	0/2	15267
Yices2	6372	6284	2176	852	8192	0/0	15684
z3	4597	7440	4826	1504	5505	0/4	18367
z3legacy	3504	6881	4081	1577	6923	891/19	16043

Table 1: Comparison among solvers on QF\_NIA

Solver	< 1s	1 to 10s	10 to 100s	100 to 600s	> 600s	unknown/unhandled	solved
CVC5	1540	1071	529	416	3391	0/0	3556
MathSat5	2995	1065	1124	1184	577	0/2	6368
Yices2	3638	2001	276	120	909	0/3	6035
z3	2840	1161	1521	754	669	0/2	6276
z3legacy	2714	1059	1619	702	851	0/2	6094

Table 2: Comparison among solvers on QF\_LIA

Solver	< 1s	1 to 10s	10 to 100s	100 to 600s	> 600s	unknown/unhandled	solved
CVC5	11	23	54	33	183	4/0	121
MathSat5	147	17	19	32	70	0/23	215
Yices2, 1000	13	21	16	12	90	0/156	62
z3	26	88	86	17	91	0/0	217
z3legacy	36	69	55	8	133	7/0	168

Table 3: Comparison among solvers on Certora Benchmarks

Disabled feature	< 1s	1 to 10s	10 to 100s	100 to 600s	> 600s	unknown/unhandled	solved
gomory-use-big-cuts	31	82	74	17	104	0/0	204
bounded-nra	30	83	81	18	96	0/0	212
branching	31	84	85	14	94	0/0	214
gomory-use-closest-int	31	83	85	16	93	0/0	215
divisions-check	29	92	77	19	91	0/0	217
enable-gcd	29	86	89	17	87	0/0	221
gomory-polarity	27	86	78	18	99	0/0	209
gomory-use-big-cuts	29	83	75	17	104	0/0	204
gomory-use-closest-int	29	86	84	15	94	0/0	214
factorization	29	86	85	18	90	0/0	218
propagate-eqs	28	90	81	20	89	0/0	219
propagate-linear	28	89	80	19	92	0/0	216
grobner	29	85	82	18	94	0/0	214
horner	29	85	81	18	95	0/0	213
incremental-linearization	30	73	70	10	125	0/0	183
monic-eq	28	89	83	18	90	0/0	218
patch-monomials	28	84	77	18	101	0/0	207
gomory-polarity	28	87	76	19	98	0/0	210
propagate-monomial-bounds	30	87	82	15	94	0/0	214

Table 4: Disabling selected features on Certora Benchmarks

Disabled feature	< 1s	1 to 10s	10 to 100s	100 to 600s	> 600s	unknown/unhandled	solved
bounded-nra	507	670	453	128	554	0/1	1758
branching	456	730	445	138	544	0/0	1769
divisions-check	462	701	470	138	541	0/1	1771
enable-gcd	461	716	449	137	549	0/1	1763
gomory-use-big-cuts	457	712	466	130	547	0/1	1765
gomory-polarity	452	724	460	135	542	0/0	1771
gomory-use-closest-int	456	709	468	139	541	0/0	1772
grobner	467	738	440	117	551	0/0	1762
horner	457	721	468	126	541	0/0	1772
incremental-linearization	402	498	286	106	1015	0/6	1292
patch-monomials	424	684	491	152	562	0/0	1751
propagate-monomial-bounds	433	697	478	156	549	0/0	1764

Table 5: Disabling selected features on a representative small subset of QF\_NIA

Disabled feature	< 1s	1 to 10s	10 to 100s	100 to 600s	> 600s	unknown/unhandled	solved
enable-gcd	2817	1151	1517	776	684	0/2	6261
gomory-use-big-cuts	2805	1164	1525	773	678	0/2	6267
gomory-use-closest-int	2815	1170	1528	758	674	0/2	6271
gomory-polarity	2850	1135	1525	750	685	0/2	6260

Table 6: Comparison of selected integer linear arithmetic features on QF\_LIA